Task 1: Rating Prediction via Prompting

Objective

To evaluate how different prompt designs affect an LLM’s ability to predict Yelp review star ratings.

Approach

- Used a subset of Yelp review data (text, stars)
- Tested three prompt strategies:
  - Simple Prompt
  - Few-shot Prompt
  - Structured Prompt
- Model responses were parsed and evaluated for:
  - Accuracy
  - JSON validity

Observations

- Results highlight prompt sensitivity and structured output challenges.
- Rate-limit constraints affected large-scale evaluation.
- Task demonstrates understanding of prompt engineering trade-offs.

Results are documented in the notebook and CSV file.

Task 2: AI Feedback Web App (Streamlit)

Objective

Build a lightweight AI-powered feedback app with an admin dashboard.

Features

User App (app.py)

- Collects user query, rating, and feedback
- Generates AI response (mock/conditional due to API quota)
- Stores feedback in data.csv

Admin Dashboard (admin.py)

- View all user feedback
- Rating-based filters
- KPIs:
  - Total feedback count
  - Average rating
  - Rating distribution visualization

Tech Stack

- Python
- Streamlit

How to Run

1. **Install Dependencies**

   pip install -r requirements.txt

2. **Run User App**

   cd task2

   streamlit run app.py

3. **Run Admin Dashboard**

   streamlit run admin.py

Notes

- AI service may display "AI service unavailable" due to free-tier rate limits.
- Logic gracefully handles API failures without breaking the app.

Submission Notes

- Focused on clarity, functionality, and clean structure
- Prioritized debuggability and readability over over-engineering
- Features and scope aligned with intern-level expectations

Author

Abhishek Shukla